

```
In [1]: import pandas as pd  
  
In [2]: import numpy as np  
  
In [3]: import matplotlib.pyplot as plt  
  
In [15]: df = pd.read_excel("Downloads/bostonhousing (1).xlsx")  
  
In [18]: df.shape  
  
Out[18]: (506, 14)
```

```
In [19]: df.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 506 entries, 0 to 505  
Data columns (total 14 columns):  
 #   Column   Non-Null Count   Dtype     
---  --  -----  -----  
 0   crim     506 non-null    float64  
 1   zn        506 non-null    float64  
 2   indus    506 non-null    float64  
 3   chas     506 non-null    int64  
 4   nox      506 non-null    float64  
 5   rm       506 non-null    float64  
 6   age      506 non-null    float64  
 7   dis       506 non-null    float64  
 8   rad      506 non-null    int64  
 9   tax      506 non-null    int64  
 10  ptratio   506 non-null    float64  
 11  b         506 non-null    float64  
 12  lstat    506 non-null    float64  
 13  medv     506 non-null    float64  
dtypes: float64(11), int64(3)  
memory usage: 55.5 KB
```

```
In [24]: boston = {  
    "data": X.values,  
    "target": y.values,  
    "feature_names": X.columns  
}  
  
print(boston["feature_names"])  
  
Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',  
       'ptratio', 'b', 'lstat'],  
      dtype='object')
```

```
In [27]: data = pd.DataFrame(df)  
data.head()
```

Out[27]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	i
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	



In [31]: `data.columns`

Out[31]: `Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax', 'ptratio', 'b', 'lstat', 'medv'], dtype='object')`

In [32]: `X = df.drop("medv", axis=1)`

In [35]: `data = X.copy()
data.columns = X.columns
data.head()`

Out[35]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33

In [34]: `data.rename(columns={'medv': 'PRICE'}, inplace=True)
data.head()`

Out[34]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33

In [36]: `data.isnull().sum()[data.isnull().sum() > 0]`

Out[36]: `Series([], dtype: int64)`

```
In [40]: data.columns = data.columns.str.strip()
```

```
In [42]: target_column = data.columns[-1]
```

```
In [43]: x = data.drop(target_column, axis=1)
y = data[target_column]
```

```
In [44]: print("Target column is:", target_column)
print(x.shape, y.shape)
```

```
Target column is: lstat
(506, 12) (506,)
```

```
In [45]: from sklearn.model_selection import train_test_split

xtrain, xtest, ytrain, ytest = train_test_split(
    x, y,
    test_size=0.2,
    random_state=0
)
```

```
In [46]: print(xtrain.shape)
print(xtest.shape)
print(ytrain.shape)
print(ytest.shape)
```

```
(404, 12)
(102, 12)
(404,)
(102,)
```

```
In [47]: from sklearn.linear_model import LinearRegression
```

```
In [48]: lm = LinearRegression()
```

```
In [49]: model = lm.fit(xtrain, ytrain)
```

```
print("Model training completed ✅")
```

```
Model training completed ✅
```

```
In [53]: from sklearn.linear_model import LinearRegression

lm = LinearRegression()
lm.fit(xtrain, ytrain)
```

```
Out[53]: ▾ LinearRegression ⓘ ⓘ
```

```
LinearRegression()
```

```
In [54]: ytrain_pred = lm.predict(xtrain)
ytest_pred = lm.predict(xtest)
```

```
In [55]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score  
  
print("Train R2:", r2_score(ytrain, ytrain_pred))  
print("Test R2:", r2_score(ytest, ytest_pred))
```

Train R2: 0.6830386021775365
Test R2: 0.4970790902412068

```
In [56]: from sklearn.metrics import mean_squared_error, r2_score
```


```
In [57]: mse_test = mean_squared_error(ytest, ytest_pred)  
print("Test MSE:", mse_test)
```

Test MSE: 21.085848548855292

```
In [58]: mse_train = mean_squared_error(ytrain, ytrain_pred)  
print("Train MSE:", mse_train)
```

Train MSE: 16.84458940970367

```
In [61]: plt.scatter(ytrain, ytrain_pred, c='blue', marker='o', label='Training data')  
plt.scatter(ytest, ytest_pred, c='lightgreen', marker='s', label='Test data')  
  
plt.xlabel('True values')  
plt.ylabel('Predicted values')  
plt.title("True value vs Predicted value")  
  
plt.legend(loc='upper left')  
plt.show()
```



In []: